An Interactive Musical Prediction System with Mixture Density Recurrent Neural Networks

Charles P. Martin
Department of Informatics and RITMO
University of Oslo, Norway
charlepm@ifi.uio.no

Jim Torresen
Department of Informatics and RITMO
University of Oslo, Norway
jimtoer@ifi.uio.no

ABSTRACT

This paper is about creating digital musical instruments where a predictive neural network model is integrated into the interactive system. Rather than predicting symbolic music (e.g., MIDI notes), we suggest that predicting future control data from the user and precise temporal information can lead to new and interesting interactive possibilities. We propose that a mixture density recurrent neural network (MDRNN) is an appropriate model for this task. The predictions can be used to fill-in control data when the user stops performing, or as a kind of filter on the user's own input. We present an interactive MDRNN prediction server that allows rapid prototyping of new NIMEs featuring predictive musical interaction by recording datasets, training MDRNN models, and experimenting with interaction modes. We illustrate our system with several example NIMEs applying this idea. Our evaluation shows that real-time predictive interaction is viable even on single-board computers and that small models are appropriate for small datasets.

Author Keywords

machine learning, recurrent neural network, mixture density network, prediction, interaction

CCS Concepts

•**Applied computing** → **Sound and music computing;** •**Computing methodologies** → **Neural networks;** •**Human-centered computing** → *Interaction paradigms;*

1. INTRODUCTION

In this paper, we consider how mixture density recurrent neural networks (MDRNNs) [1, 7] can be applied to real-time gestural prediction in new interfaces for musical expression (NIMEs) and present an interactive system for training and applying MDRNNs to a broad range of NIMEs. Research applying deep learning to music *generation* is rapidly appearing, but few of these systems have been applied in the service of real-time musical *performance.* We feel that deep neural networks (DNNs) can extend creative possibilities for NIME performers and designers; however, these users need DNN models that are appropriate for typical NIME-data, and better tools to allow rapid-prototyping and creative exploration with these models.

Figure 1: Our interactive musical prediction system allows NIMEs to be integrated with a mixture density recurrent neural network (MDRNN) that can support or accompany a human performance.

Present work in musical AI is usually focussed on high-level symbolic music; however, most NIME control data is better represented as low-level continuous sensor values. We propose to use MDRNNs to model this data, including time-deltas between each reading. This approach has the advantage of modelling music at the embodied [12] control level; such models imitate performing on instruments, not composing music. Another advantage is in representing rhythms absolutely—as a sequence of real-valued time-deltas—rather than being limited to a sixteenth-note grid. MDRNNs have previously been applied to control data in sketching [8] and handwriting [7], both creative tasks.

We have developed an interactive musical prediction system (IMPS), to accelerate development of musical MDRNN models that predict musical control data in real-time during performance (see Figure 1). This system assists with data collection, model training, and real-time inference. IMPS connects to typical NIME software and devices via a simple open sound control (OSC) interface and can apply an MDRNN through a number of new predictive interaction paradigms to accompany or support performers. This tool provides a new solution for MDRNN-based NIME development for artists and computer musicians. We imagine that artists could train MDRNN models on small datasets of interaction data, tailored to commercial or DIY interfaces applied in their practice. While small models may not represent all possible musical interactions, they might perform well enough to imitate aspects of an artist's style.

The main contributions of this research are the novel IMPS system that assists with data-collection, training, and real-time application of MDRNNs. We discuss the design of this system, and in particular the advantages of an MDRNN model over other popular DNNs. Our evaluation is focussed on the training and application of this system in real-time performance. We describe a number of example NIMEs developed using these tools, and show that incorporation of real-time MDRNN predictions is feasible on even single-board computers such as the Raspberry Pi.
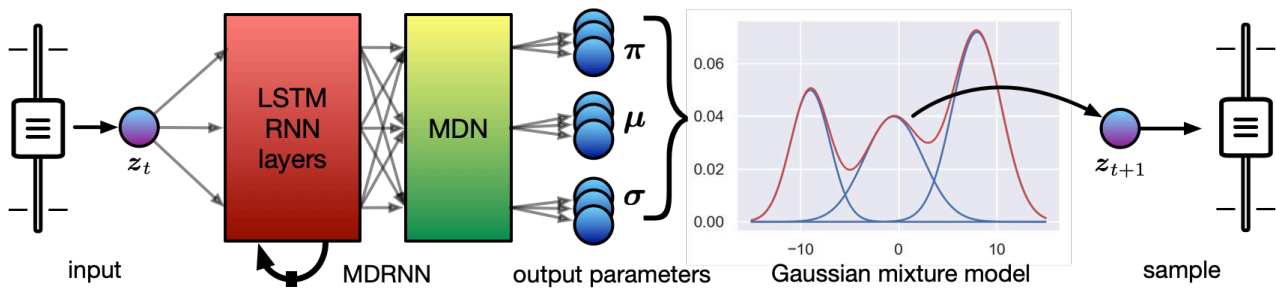
Figure 2: An MDRNN transforms the output of an RNN to form the parameters of a Gaussian mixture model. Our system allows a configurable MDRNN to be applied to predicting musical control data, such as mixer sensors or touchscreen data.

2. MACHINE LEARNING IN NIMES

Machine learning (ML) has been a part of NIME designs since at least the early 1990s [11]. More recently, the adoption of ML algorithms into NIMEs has accelerated due to frameworks such as Wekinator [4], SARC EyesWeb Catalogue [5] or ml.lib [2], that allow training and application of ML algorithms through a GUI or computer music tools such as Max or Pd.

ML has typically been applied in NIMEs for two tasks: mapping from gesture to high-level synthesiser commands, and modelling of ongoing musical processes [14]. On the mapping side, classification algorithms and shallow NNs have been useful for mapping multidimensional control data such as wind instrument keys [18] to musical notes or high level gestures. Modelling algorithms such as Markov models [17] or factor oracles [13], have been used to develop systems that learn to imitate musical styles, or interact with an improvising musician. These modelling systems usually operate on high-level musical notes or acoustic phrases, and as a result have not been as applicable to data from the multiple continuous controllers of many NIMEs.

2.1 Deep Neural Networks for Music

In recent years much research on musical applications of DNNs has appeared (see [19] for a survey). Artificial neural networks (ANNs) process data by applying simple mathematical transformations grouped into units, inspired by neurons in the brain. The strength of connections between units, how much data is allowed to move between them, are the parameters of an ANN. By arranging units deeply, into layers, ANNs be taught to accomplish high-level tasks such as recognising images and generating musical melodies.

For music applications, multi-layered recurrent neural networks (RNNs) with long short-term memory (LSTM) units are often used to learn sequences of symbolic music data. In RNNs, data can be stored in between computations and fed back as extra input. This means that RNNs can learn to understand temporal relationships between sequences of data that are applied to their inputs [6]. LSTM units are RNN components that contain an additional memory state and four internal gates that control how data is stored, released, retained and "forgotten" from the memory state.

Typical LSTM music models generate music by predicting one discrete note symbol at a time using MIDI- [10] or text-inspired representations [19]. Rather than directly predicting a note, the outputs of these networks form a categorical (softmax) probability distribution from which one discrete option can be sampled. The temperature (or diversity) of this distribution can be adjusted to favour the highest scoring option, to boost the chance of sampling less likely notes.

Symbolic LSTM RNNs can be used to model discrete musical data (e.g., notes from a piano keyboard on a 16th-note rhythmic grid), but a different architecture is required to directly predict *continuous-valued* musical data such as input from touch sensors. To model touchscreen data, Hantrakul [9] used a RNN without a stochastically sampled output. Mixture density networks (MDN) have been applied to touchscreen performances [15], which preserves stochastic sampling. That system added direct prediction of rhythm enabling pauses of arbitrary length, taps, as well as continual swirls. As will be further explained below, MDNs are highly applicable to many types of continuous control data, not just touchscreens.

While tools to define and train DNNs (e.g., Keras and TensorFlow) have become more accessible, the lack of a NIME-focussed toolkit has held back the application of DNNs within an interactive music context. This research attempts to fill this gap with an MDRNN prediction system that can be explored without additional programming.

2.2 Mixture Density RNNs

A mixture density network (MDN) transforms the outputs of a neural network into the parameters of a Gaussian mixture model (GMM) [1], as shown in Figure 2. Such a model "mixes" a number of Gaussian (or normal) distributions (in blue) with weights corresponding to the likelihood of each components, to form a more complex distribution (in red). This allows the model to represent phenomena that appear to be drawn from multiple distributions. As shown in Figure 2, the MDN's output parameters consist of centres (μ) and scales (σ) for each component distribution, and a weight (π) for each component. An MDN can be applied to the outputs of an RNN forming an MDRNN that can be trained on sequential data such as 2D pen movements for handwriting [7] and sketches [8], as well as musical touchscreen data [15].

In a creative process such as musical improvisation, multiple choices for the next note to perform or action to take could be artistically valid. This suggests that some kind of multi-modal distribution, such as a GMM, would be appropriate to accurately model such a process. An MDRNN thus forms a useful network for regression problems involving multiple correct answers, or that require a certain amount of stochasticity when sampling, as in creative tasks.

One of the complexities of an MDN is the error (loss) function, used for training, which is derived from the probability density function of the mixture model. While this is straightforward for a 1-dimensional [1] or 2D [7] case, mixtures of higher-dimensional Gaussian distributions are more complex. Our system, outlined below, makes use of recent features of TensorFlow to generate a tractable loss function for arbitrary many dimensions of data.

3. SYSTEM DESIGN

Our predictive musical interaction system consists of two components: a Keras implementation of an MDRNN with an MDN layer that extends to sequential data of arbitrary dimension and a set of Python applications to facilitate data collection, model training, and real-time interactive prediction. The system is open source[1] and operated through a command line interface that enables these tasks to be performed within the context of a NIME prototyping and performance process. This system is designed to allow artists, as well as machine learning researchers, to apply MDRNNs to creative work and other applications.

3.1 MDN Layer

Our MDN layer allows the selection of the number of mixture components and the dimension of the data to be modelled. Each mixture component is a multivariate Gaussian distribution limited to a diagonal covariance matrix. So for K components and dimension N, for one prediction, the MDN generates K mixing coefficients ($\boldsymbol{\pi}$), $K \times N$ means ($\boldsymbol{\mu}$) and $K \times N$ scales ($\boldsymbol{\sigma}$, diagonals of the covariance matrix). This is illustrated for $N = 1$, $K = 3$ in Figure 2 with an example of the resulting mixture distribution.

The limited covariance matrix means that the loss function can be easily calculated using TensorFlow Probability's `Mixture`, `Categorical`, and `MultivariateNormalDiag` functions. The loss function for an MDN relies on the number of mixture components as well of the dimension of each component, so this function is generated on demand for these parameters.

3.2 MDRNN

We provide an abstraction, `PredictiveMusicMDRNN`, for generating an appropriate MDRNN for learning musical sequences. This is constructed using multiple layers of LSTM units followed by an MDN layer. The number of layers, LSTM units, mixture components, and distribution dimensions are configurable hyperparameters, but we suggest that 2-layer networks with 5 mixture components are used. The abstraction assumes that the first dimension of the data will be used for the time since the previous sample (dt) and will thus be positive and nonzero. The other dimensions of the data are assumed to be between 0 and 1. This means that to model data with 2 continuous control variables (x_1 and x_2), a 3D MDRNN is required (dt, x_1, x_2). Our abstraction can construct networks for training (unrolled to the correct training sequence length) and for inference (with sequence length set to 1, and LSTM state stored in between calculations), and provides a function to make inferences one at a time as input data becomes available.

We provide a high-level control of the size of the MDRNN which allows the user to trade potential learning capacity for speed of training and inference. Four presets are provided that change the number of LSTM units as follows: s–64, m–128, l–256, xl–512. The 's' network has around 50K parameters while the 'xl' network has 3M. Further customisation to hyperparameters is easily made through a command line interface.

3.3 Predictive Musical Interaction Controller

The main IMPS application receives interaction messages from a control interface over OSC (e.g., `/interface`, x_1, x_2, . . .). All interaction messages are automatically logged to CSV files in order to build up a training dataset.

IMPS' MDRNN module can accept interface messages as input, performing inference and sampling to produce pre-

Figure 3: The EMPI self-contained NIME (enclosure open) represents predictions physically and sonically. It uses a 2D MDRNN (dt, input value).

dictions of the user's next control interaction (as illustrated in Figure 2). Information from the MDRNN's input conditions the LSTM layers' internal states, which means that predictions are influenced by previous interactions. The MDRNN can also perform predictions on top of predictions by connecting output to input, allowing it to generate new control signals based on previous experiences. Predictions are stored internally and sent back to the interface (with the address `/prediction`), at the time indicated by the dt values output by the MDRNN. This allows the MDRNN to perform arbitrarily timed rests and rhythms.

The input and output of the MDRNN can be connected in different ways to messages from the interface and the outgoing stream of predicted control values. Given that messages may arrive faster than MDRNN predictions are made, IMPS uses concurrent queues to keep up as much as possible. We have defined four interaction modes in IMPS to explore these configurations in performance:

No Predictions. Interface messages condition the MDRNN memory, but the output predictions are ignored.
Filter. Interface messages condition the MDRNN memory; output predictions are played as quickly as possible.
Call-and-Response. The MDRNN is conditioned when the user plays; if the user stops, the MDRNN generates new control signals until they interact again.
Battle. The user and MDRNN play together, with the MDRNN continually generating control signals regardless of the user's actions.

These predictive interactions are a starting point for further research. In particular, they demonstrate that our system is capable of both mapping (through the filter interaction), and modelling (through call-and-response).

4. EXAMPLE APPLICATIONS

In this section we present several NIMEs that integrate our MDRNN predictive model to demonstrate its flexibility for a variety of interactive music modalities. These systems allow the output from the MDRNN to be sonified as well as represented visually or physically. The systems have differing degrees of freedom, from one single controller, to an eight-knob interface. These systems are illustrated in our video abstract[2].

4.1 EMPI

The Embodied Musical Predictive Interface (EMPI) is a self-contained NIME with a single dimension of continuous input and output through two physical levers, a Raspberry Pi, touchscreen, and speaker. EMPI was designed to explore the simplest predictive musical interactions: where one dimension of input is modelled along with time. The MDRNN model was trained on a 10-minute human performance with

---

[1]IMPS Code: https://doi.org/10.5281/zenodo.2580175

[2]Video Abstract: https://doi.org/10.5281/zenodo.2597494

Figure 4: The LightPad is a pressure-sensitive multitouch controller with an LED display. It can be used with 4D MDRNN model (dt, x, y, pressure)

Figure 5: We used the X-Touch Mini to control an additive synthesis instrument along with a 9D MDRNN model (dt, knobs 1–8).


Figure 6: Training and validation loss for a small LightPad dataset (12K interactions) for different models (lower is better). The 64-unit model had the best performance on the validation data.


Figure 7: Training and validation loss for a 100K LightPad interaction dataset. The 64- and 128-unit models had the best validation loss.

the input lever. The EMPI demonstrates that even a Raspberry Pi can be used for predictions from a small MDRNN in a real-time situation. In performance, the predictions are sonified through the EMPI's speaker as well as physically represented through the servo-controlled lever, this imbues the NIME with a sense of independent agency.

4.2 LightPad Block

The Roli LightPad Block (see Figure 4) is a small multitouch controller with a flexible pressure sensitive surface and an LED display. The LightPad's expressive inputs and the ability to visually represent these inputs through the display make it ideal for experimenting with predictive interaction. We developed a Pd patch that communicates between our predictive interaction system (OSC) and the LightPad (MIDI over Bluetooth). This allows MIDI messages from the LightPad, and predictions from the MDRNN to be mapped to software synthesisers. The LightPad's LED display is usually used to illuminate the user's touches as well as the state of music software. We feed MDRNN predictions back to the LightPad which are displayed in a different colour. This allows an intuitive view of the user's actions and predictive responses during performance.

The LightPad data was applied to a 4D MDRNN (dt, x-position, y-position, pressure). Hantrakul previously presented an RNN LightPad controller [9], however that system used a deterministic RNN without the ability to predict *when* to play the LightPad. Unlike that system, IMPS can stop swirling and predicts when and where to start a new touch, a significant advantage allowing the performance of novel rhythms as well as gestures.

4.3 X-Touch Mini

Mixing control surfaces with multitudes of faders and knobs are common in recording studios as well as in NIME research. Interfaces with illuminated controls, such as the Behringer X-Touch Mini, can represent previously recorded data as it is played back (see Figure 5). We used the X-Touch Mini to control the parameters of an additive synthesis instrument, and recorded 10-minutes of performance with this instrument

that comprised 12,000 interactions. This data was used to train a 9D MDRNN model (dt, knobs 1–8). In performance, predictive output can be sent to a synthesis routine as well as represented visually with the illuminated controls. ML fader interfaces have previously been demonstrated without a neural-network powered controller [20]. Unlike Tahiroglu et al's research, our system can interpolate between the many training examples and come up with novel gestures.

5. EVALUATIONS AND DISCUSSION

Many aspects of predictive music systems can be evaluated to determine their place in a NIME context; however, a key question for NIME designers and performers wishing to apply IMPS would be whether training is fast enough to keep up with a rapid-prototyping process and predictions are fast enough to keep up with real-time performance. In this section we evaluate training performance on small datasets, prediction speed, and strategies for sampling useful predictions.

We note here that training and inference are very different operations for the MDRNN. In short, while inference only involves computing the output from a single timestep, training involves unrolling the MDRNN to multiple (our default is 50) timesteps to calculate the gradients of the loss function, this is computed in parallel for each member of a training batch (our default batch size is 64), thus one training step requires much more computation than one inference calculation.

5.1 Training

Training, as outlined above, is much slower than inference for DNNs. While training very large DNNs is slow even on powerful GPUs, the small MDRNNs discussed here can be trained reasonably quickly on a small dataset, and on a normal computer. The main hyperparameter to consider

Figure 8: Time (mean and SD) taken per prediction for different MDN output dimensions (LSTM units = 64) showing little effect on computation speed.

Figure 9: Time (mean and SD) taken per prediction for different RNN layer sizes. The time is <10ms for 64-unit networks on all systems.

here is the size of the MDRNN compared to the training dataset; larger MDRNNs have more learning capacity but are slower to train and may not learn more from a smaller dataset. By way of example, we consider the LightPad model discussed in Section 4.2. We trained LightPad models on two datasets: a very small corpus of 12K touch interactions (15 minutes of performance), and a larger corpus of 100K interactions (2 hours). The small corpus takes 70 minutes to train for 100 epochs (exposures to all examples) on a MacBook Pro.

Would it help in this case to use a larger network? We trained 64-, 128-, 256-, and 512-unit networks on each dataset, and used the loss values to investigate learning performance as shown in Figure 6 and 7. The training loss is the average loss over each batch of training examples, and the validation loss is calculated on a set (10% in our case) of examples that are held out to evaluate training after each epoch. Figure 6 shows that all but the 64-unit network had poor performance against the validation set on the 12K dataset indicating that the larger networks had overfit to the training data. For the 100K dataset, Figure 7 shows that both the 64- and 128-unit networks had reasonable validation performance. It is likely that the smaller networks are most useful for datasets of this size. Early stopping (an option in IMPS) would have saved much time here by stopping training after the validation loss failed to improve.

While 70 minutes of training time is more than the mere seconds required to train Wekinator models [4], we argue that it does not preclude rapid NIME-prototyping. Similarly to Wekinator, an MDRNN with our system can be improved by subsequent retraining. Our predictive interaction system logs all interactions to facilitate the accumulation of datasets as a NIME is refined. In our process, models are continually retrained during NIME creation as more data becomes available through experimentation and performance.

5.2 Prediction Speed

We benchmarked prediction speed from our MDRNN on four computer systems: a MacBook Air (*mba*, Intel i5 1.8GHz), MacBook Pro (*mbp*, Intel i7 2.6GHz), desktop PC (*gpu*, NVidia GeForce GTX 1080TI), and a Raspberry Pi 3 B+ (*rpi*, Broadcom BCM2837B0 1.4GHz). The desktop PC computed MDRNN predictions on the GPU and the other systems used their CPUs. We examined two aspects of the MDRNN configuration—the input/output dimension of the network (between 2 and 9), and the number of LSTM units in each RNN layer (64, 128, 256, 512). 100 predictions were performed with each configuration on each system with the first of each test discarded due to time taken for setup overhead.

The prediction time for a 64-unit per layer MDRNN with different output dimension sizes is shown in Figure 8. This shows that changing the output dimension from 2 up to 9 has little impact on computation speed for each computer system. The final MDN layer uses few parameters in comparison to the RNN layers in the network, so its size has little impact on computation time. A practical consequence is that it is feasible to experiment with MDRNN prediction of complex NIMEs with many dimensions of real-valued input.

The number of LSTM units in each RNN layer has a much more significant impact on the speed of prediction (Figure 9). The *mbp* achieved the best speed for the 64 unit network with a mean of 1.85ms per prediction and a mean of only 4.51ms on the 512 unit network. This means that MDRNN predictions could fit within typical expectations for NIME latency of around 10ms [16] on this system. For *gpu*, prediction speed barely changed between 64- (mean=2.85ms) and 512-unit (mean=2.96ms) networks due to the advantages of GPU-acceleration for DNN calculations. The *rpi* saw the greatest changes in speed, from a mean of 6.69ms on the 64-unit network to 37.64ms on the 512-unit network. While smaller networks are feasible on the *rpi*, larger networks may be too slow in practice.

5.3 Sampling

The output from an MDN is the set of parameters for the GMM: the weights for each mixture component ($\boldsymbol{\pi}$), and the means and scales for each multivariate Gaussian ($\boldsymbol{\mu}$ and $\boldsymbol{\sigma}$). Predictions must be generated from these parameters by sampling from the categorical distribution formed by $\boldsymbol{\pi}$, and then sampling from the Gaussian distribution chosen by that outcome.

The concept of adjusting the temperature of a categorical distribution will be familiar to those who have used RNNs to learn creative sequences such as text or symbolic music. The MDN's categorical distribution can be adjusted in the same way with very low temperature values favouring the maximum value in the distribution, and high values producing a more uniform distribution. This adjustment could be called "π-temperature". The temperature of the Gaussian distributions can also be adjusted by scaling the variances $\boldsymbol{\sigma}$. High values result in a wider spread of predictions, and low values are closer to the selected mean. We call this "σ-temperature".

In IMPS, we have found adjusting the σ- and π-temperature to be very important for making useful predictions. Figure 5.3 shows touchscreen performances on a 3D (dt, x, y) network generated at different temperatures. The MDRNN can end up producing high $\boldsymbol{\sigma}$s, resulting in jagged output, but by reducing the σ-temperature to close to zero, we can
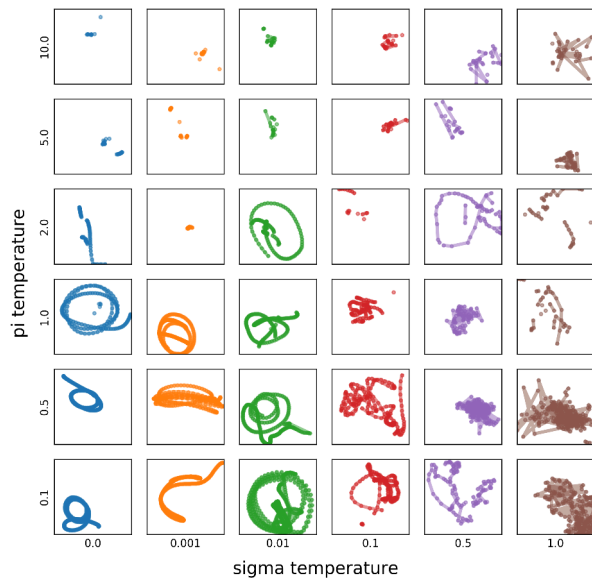
Figure 10: Performances from a 3D (dt, x, y) MDRNN at different sampling temperatures.

generate smooth paths. The π-temperature can control the appearance of different gestures to some extent. At low values, an MDRNN will have trouble changing modes, endlessly swirling. At very high values, it taps without completing any significant swipes. Exploring how π and σ sampling temperature can be explored in our predictive interaction systems is a topic of our future research.

6. CONCLUSION AND FUTURE WORK

We have introduced IMPS, an Interactive Musical Prediction System, that allows an MDRNN model to be applied in any NIME with continuous control signals. This system has been demonstrated through predictive interactions with three interfaces: the custom built EMPI, commercial LightPad, and X-Touch Mini. Our evaluation has shown that IMPS is viable for use on normal laptops and even single-board computers, and provided insight into how the MDRNN behaves under different training and sampling parameters. It is promising that small, easily trained models seem sufficient to learn smaller datasets; however, future studies could seek to understand what aspects of a control style are actually learned and how IMPS might work within a design process.

We think that predictive interaction could be applied widely in computer music software such as DAWs, synthesis environments and physical hardware controllers. The success of tools such as Wekinator [3], and interest in Google's Magenta project suggest that artists see the value of applying ML in their work. The flexibility of IMPS and our MDRNN design could be ideal for providing predictive interaction possibilities to these users.

Acknowledgements

7. REFERENCES

[1] C. M. Bishop. Mixture density networks. Technical Report NCRG/97/004, Neural Computing Research Group, Aston University, 1994.

[2] J. Bullock and A. Momeni. ml.lib: Robust, cross-platform, open-source machine learning for Max and Pure Data. In *Proc. NIME*, pages 265–270, 2015.

[3] R. Fiebrink. Machine learning as meta-instrument: Human-machine partnerships shaping expressive instrumental creation. In *Musical Instruments in the 21st Century: Identities, Configurations, Practices*, pages 137–151. Springer, 2017.

[4] R. Fiebrink, D. Trueman, and P. R. Cook. A meta-instrument for interactive, on-the-fly machine learning. In *Proc. NIME*, pages 280–285, 2009.

[5] N. Gillian, R. B. Knapp, and S. O'Modhrain. A machine learning toolbox for musician computer interaction. In *Proc. NIME*, pages 343–348, 2011.

[6] I. Goodfellow, Y. Bengio, and A. Courville. *Deep Learning*. MIT Press, 2016.

[7] A. Graves. Generating sequences with recurrent neural networks. *ArXiv e-prints*, 2013. `arXiv:1308.0850`.

[8] D. Ha and D. Eck. A neural representation of sketch drawings. *ArXiv e-prints*, 2017. `arXiv:1704.03477`.

[9] L. Hantrakul. GestureRNN: A neural gesture system for the Roli Lightpad Block. In *Proc. NIME*, pages 132–137, 2018.

[10] N. Jaques, S. Gu, D. Bahdanau, J. M. Hernández-Lobato, R. E. Turner, and D. Eck. Sequence tutor: Conservative fine-tuning of sequence generation models with KL-control. In *Proc. ICML*, pages 1645–1654, 2017.

[11] M. Lee, A. Freed, and D. Wessel. Real-time neural network processing of gestural and acoustic signals. In *Proc. ICMC*, 1991.

[12] M. Leman, P.-J. Maes, L. Nijs, and E. Van Dyck. What is embodied music cognition? In R. Bader, editor, *Springer Handbook of Systematic Musicology*, pages 747–760. Springer, Berlin, Heidelberg, 2018.

[13] B. Lévy, G. Bloch, and G. Assayag. Omaxist dialectics. In *Proc. NIME*, pages 137–140, 2012.

[14] C. P. Martin, K. Olav Ellefsen, and J. Torresen. Deep Predictive Models in Interactive Music. *arXiv e-prints*, 2018. `arXiv:1801.10492`.

[15] C. P. Martin and J. Torresen. Robojam: A musical mixture density network for collaborative touchscreen interaction. In A. Liapis, J. J. Romero Cardalda, and A. Ekárt, editors, *Computational Intelligence in Music, Sound, Art and Design*, pages 161–176, 2018.

[16] A. McPherson, R. Jack, and G. Moro. Action-sound latency: Are our tools fast enough? In *Proc. NIME*, pages 20–25, 2016.

[17] F. Pachet. The continuator: Musical interaction with style. *J. New Music Res.*, 32(3):333–341, 2003. `doi:10.1076/jnmr.32.3.333.16861`.

[18] J. Snyder and D. Ryan. The birl: An electronic wind instrument based on an artificial neural network parameter mapping structure. In *Proc. NIME*, pages 585–588, 2014.

[19] B. L. Sturm and O. Ben-Tal. Taking the models back to music practice: Evaluating generative transcription models built using deep learning. *J. Creative Music Sys.*, 2(1), 2017.

[20] K. Tahiroglu, T. Svedström, and V. Wikström. Musical engagement that is predicated on intentional activity of the performer with NOISA instruments. In *Proc. NIME*, pages 132–135, 2015.